

1 OpenType Font Variations Common Table

2 Formats

3 OpenType Font Variations allow continuous variation along one or more design axes,
4 such as weight or width. An overview of OpenType Font Variations and a specification of
5 the algorithm for interpolating variation instance values is provided in the chapter,
6 [OpenType Font Variations Overview](#); that chapter should be read first. This chapter
7 documents the formats for variation data that are used in various font tables, such as
8 the 'gvar' or [MVAR](#) tables. The data stored using the formats described in this chapter
9 are processed as described in detail in the Overview chapter; additional, higher-level
10 information on processing is provided here.

11 Overview

12 A font has many different data items found in several different font tables that provide
13 values that are specific to a particular font face. Examples include glyph-specific values,
14 such as the positions of glyph outline points and glyph advance widths, and face-wide
15 values, such as a sub-family name, a weight class, or ascender and descender values. In a
16 variable font, most or all of these values may need to vary for different variation
17 instances. When an application selects a variation instance within the font's variation
18 space, new values for such items appropriate to that instance need to be derived. This is
19 done using delta adjustment values.

20 For example, the OS/2 table of a font may provide a default `sxHeight` value of 970. The
21 [MVAR](#) table might provide a delta value of +50 that is used for weight-axis values from
22 the default to the heaviest-supported weight. For a particular instance, the interpolation
23 process might scale that delta with a scalar co-efficient of 0.4, deriving an instance
24 `sxHeight` value of 990.

25 These concepts and the interpolation algorithm for deriving instance values are
26 described in detail in the chapter, [OpenType Font Variations Overview](#).

27 The variation data for a font consists of a number of delta adjustment values. Each
28 individual delta applies to a particular, target data item — for instance, the X coordinate
29 of a point in a glyph outline, or the font's `sTypoAscender` — and is also associated with
30 a specific region within the font's design variation space over which it is applicable.
31 Thus, a given delta is logically keyed by the target data item and the applicable region.

32 A variable font includes many deltas. At the highest level, deltas are organized into
33 collections for different target item sets:

- 34 • Deltas for positions of points of a 'glyf' table are stored in a 'gvar' table.
- 35 • Deltas for positions of points of a CFF2 table are stored within the CFF2 table.
- 36 • Deltas for CVT values in the 'cvt' table are stored in a 'cvar' table.
- 37 • Deltas for glyph metrics in an 'hmtx' table are stored in an HVAR table; and deltas
38 for glyph metrics in a 'vmtx' or VORG table are stored in a VVAR table.
- 39 • Deltas for anchor positions in GPOS lookups and other items used in GDEF, GPOS
40 or JSTF tables are stored within variation data contained in the GDEF table.
- 41 • Deltas for font-wide metrics and other items from the OS/2, 'hhea', 'gasp', 'post'
42 or 'vhea' tables are stored in an MVAR table.
- 43 • Deltas for values in other tables are stored in the respective table: deltas for
44 baseline metrics in the BASE table and for various items in the COLR table are
45 stored in each table.

46 In a variable font, the largest group of deltas are for the positions of glyph outline
47 points. For TrueType outlines in a 'glyf' table, the deltas are stored within the 'gvar'
48 table, with a second level of organization grouping deltas by glyph ID. See the 'gvar'
49 table specification for details.

50 Below these higher levels of organization, most variation data is organized in one of two
51 ways. (Variation data for CFF 2 outlines is a partial exception — see below.)

- 52 • Organize sets of deltas for several target items into groupings by the variation-
53 space region over which they apply. Since regions are defined using n-tuples (or
54 "tuples"), such data sets will be referred to as *tuple variation stores*.
- 55 • Organize sets of deltas associated with different regions into groupings by the
56 target items to which they apply. Such data sets will be referred to as *item*
57 *variation stores*.

58 The two formats have different ways of representing n-tuples that define regions of
59 applicability, and different ways of associating deltas with target font-data items. The
60 tuple variation store format is optimized for compact representation of glyph outline
61 variation data that is all processed for a given variation instance. The item variation store
62 format, on the other hand, is designed to allow direct access to variation data for
63 arbitrary target items, allowing more efficient processing in contexts that do not require
64 interpolated values for all items to be computed. (Additional details are provided
65 below.) The 'gvar' and 'cvar' table use the tuple variation store format, while variation
66 data in most other tables, including the MVAR, HVAR and GDEF tables, use item
67 variation store formats.

68 Variation data for CFF 2 outlines are handled slightly differently than other cases. The
69 deltas for glyph outline descriptions are interleaved directly within the outline
70 descriptions in the [Compact Font Format 2 \(CFF2\)](#) table. But the sets of regions that are
71 associated with the delta sets are defined in an item variation store, contained as a
72 subtable within the CFF2 table.

73 Tuple Variation Store

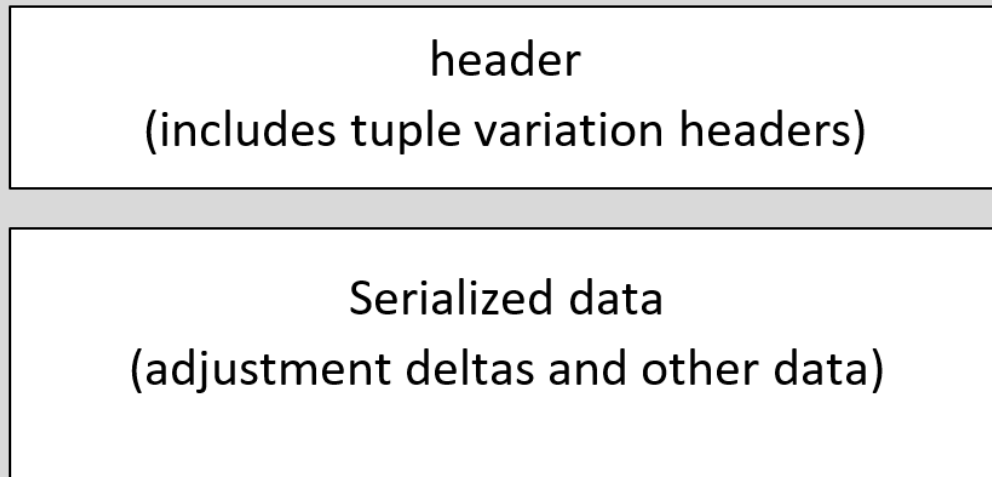
74 Tuple variation stores are used in the 'gvar' and 'cvar' tables, and organize sets of
75 variation data into groupings, each of which is associated with a region of applicability
76 within the variation space. Within the 'gvar' table, there is a separate variation store for
77 each glyph. Within the 'cvar' table, there is one variation store providing variations for all
78 CVT values.

79 There is a minor difference in the top-level structure of the store in these two contexts.
80 Within the 'cvar' table, it is the entire 'cvar' table that comprises the specific variation
81 store format, with a header that begins with major/minor version fields. The specific
82 variation store format for glyph-specific data within the 'gvar' table is the
83 *GlyphVariationData* table (one per glyph ID), which does not include any version fields.
84 In other respects, the 'cvar' table and *GlyphVariationData* table formats are the same.
85 There is also a minor difference in certain data that can occur in a *GlyphVariationData*
86 table versus a 'cvar' table. Differences between the 'gvar' and 'cvar' tables will be
87 summarized later in this section.

88 In terms of logical information content, the *GlyphVariationData* and 'cvar' tables consist
89 of a set of logical, tuple variation data tables, each for a different region of the variation
90 space. In physical layout, however, the logical tuple variation tables are divided into
91 separate parts that get stored separately: a header portion, and a serialized-data
92 portion.

93 In terms of overall structure, the *GlyphVariationData* table and the 'cvar' table each
94 begin with a header, which is followed by serialized data. The header includes an array
95 with the tuple variation headers. The serialized data include deltas and other data that
96 will be explained below.

glyphVariationData table / 'cvar' table



97
98 Figure: High-level organization of tuple variation stores

99 Tuple Records

100 The tuple variation store formats reference regions within the font's variation space
101 using tuple records. These references identify positions in terms of normalized
102 coordinates, which use F2DOT14 values.

103 *Tuple record (F2DOT14):*

Type	Name	Description
F2DOT14	coordinates[axisCount]	Coordinate array specifying a position within the font's variation space. The number of elements must match the axisCount specified in the 'fvar' table.

104 Tuple Variation Store Header

105 The two variants of a tuple variation store header, the GlyphVariationData table header
106 and the 'cvar' header, are only slightly different. The formats of each are as follows:

107 *GlyphVariationData header:*

Type	Name	Description
uint16	tupleVariationCount	A packed field. The high 4 bits are flags (see below), and the low 12 bits are the number of tuple variation tables for this glyph. The count can be any number between 1 and 4095.
Offset16	dataOffset	Offset from the start of the GlyphVariationData table to the serialized data.
TupleVariationHeader	tupleVariationHeaders[tupleVariationCount]	Array of tuple variation headers.

108 'cvar' table header:

Type	Name	Description
uint16	majorVersion	Major version number of the 'cvar' table — set to 1.
uint16	minorVersion	Minor version number of the 'cvar' table — set to 0.
uint16	tupleVariationCount	A packed field. The high 4 bits are flags (see below), and the low 12 bits are the number of tuple variation tables. The count can be

		any number between 1 and 4095.
Offset16	dataOffset	Offset from the start of the 'cvar' table to the serialized data.
TupleVariationHeader	tupleVariationHeaders[tupleVariationCount]	Array of tuple variation headers.

109 The tupleVariationCount field contains a packed value that includes flags and the
110 number of logical tuple variation tables — which is also the number of physical tuple
111 variation headers. The format of the tupleVariationCount value is as follows:

Mask	Name	Description
0x8000	SHARED_POINT_NUMBERS	Flag indicating that some or all tuple variation tables reference a shared set of “point” numbers. These shared numbers are represented as packed point number data at the start of the serialized data.
0x7000	Reserved	Reserved for future use — set to 0.
0x0FFF	COUNT_MASK	Mask for the low bits to give the number of tuple variation tables.

112 If the SHARED_POINT_NUMBERS flag is set, then the serialized data following the
113 header begins with packed “point” number data. In the context of a GlyphVariationData
114 table within the 'gvar' table, these identify outline point numbers for which deltas are
115 explicitly provided. In the context of the 'cvar' table, these are interpreted as CVT indices
116 rather than point indices. The format of packed point number data is described below.

117 TupleVariationHeader

118 The GlyphVariationData and 'cvar' header formats include an array of tuple variation
119 headers. The TupleVariationHeader format is as follows.

120 *TupleVariationHeader:*

Type	Name	Description
uint16	variationDataSize	The size in bytes of the serialized data for this tuple variation table.
uint16	tupleIndex	A packed field. The high 4 bits are flags (see below). The low 12 bits are an index into a shared tuple records array.
Tuple	peakTuple	Peak tuple record for this tuple variation table — optional, determined by flags in the tupleIndex value. Note that this must always be included in the 'cvar' table.
Tuple	intermediateStartTuple	Intermediate start tuple record for this tuple variation table — optional, determined by flags in the tupleIndex value.
Tuple	intermediateEndTuple	Intermediate end tuple record for this tuple variation table — optional, determined by flags in the tupleIndex value.

121 Note that the size of the TupleVariationHeader is variable, depending on whether peak
122 or intermediate tuple records are included. (See below for more information.)

123 The variationDataSize value indicates the size of serialized data for the given tuple
124 variation table that is contained in the serialized data. It does not include the size of the
125 TupleVariationHeader.

126 Every tuple variation table has an associated peak tuple record. Most tuple variation
127 tables use non-intermediate regions, and so require only the peak tuple record to define
128 the region. In the 'cvar' table, there is only one variation store, and so any given region
129 will only need to be referenced once. Within the 'gvar' table, however, there is a
130 GlyphVariationData table for each glyph ID, and so any region may be referenced
131 numerous times; in fact, most regions will be referenced within the GlyphVariationData
132 tables for most glyphs. To provide a more efficient representation, the tuple variation
133 store formats allow for an array of tuple records, stored outside the tuple variation store
134 structures, that can be shared across many tuple variation stores. This is used only within
135 the 'gvar' table; it is not needed or supported in the 'cvar' table. The formats alternately
136 allow for a peak tuple record that is non-shared, specific to the given tuple variation
137 table, to be embedded directly within a TupleVariationHeader. This is optional within the
138 'gvar' table, but required in the 'cvar' table, which does not use shared peak tuple
139 records.

140 See the 'gvar' chapter for details on the representation of shared tuple records within
141 that table.

142 The tupleIndex field contains a packed value that includes flags and an index into a
143 shared tuple records array (not used in the 'cvar' table). The format of the tupleIndex
144 field is as follows.

145 *tupleIndex format:*

Mask	Name	Description
0x8000	EMBEDDED_PEAK_TUPLE	Flag indicating that this tuple variation header includes an embedded peak tuple record, immediately after the tupleIndex field. If set, the low 12 bits of the tupleIndex value are ignored. Note that this must always be set within the 'cvar' table.
0x4000	INTERMEDIATE_REGION	Flag indicating that this tuple variation table applies to an intermediate region within the variation space. If set, the header includes the two intermediate-region, start and end tuple records, immediately after the peak tuple record (if present).
0x2000	PRIVATE_POINT_NUMBERS	Flag indicating that the serialized data for this tuple variation table includes packed "point" number data. If set, this tuple variation table uses that number data; if clear, this tuple variation table uses shared number data found at the start of the serialized data for this glyph variation data or 'cvar' table.
0x1000	Reserved	Reserved for future use — set to 0.
0x0FFF	TUPLE_INDEX_MASK	Mask for the low 12 bits to give the shared tuple records index.

146 Note that the intermediateRegion flag is independent of the embeddedPeakTuple flag
147 or the shared tuple records index. Every tuple variation table has a peak n-tuple
148 indicated either by an embedded tuple record (always true in the 'cvar' table) or by an
149 index into a shared tuple records array (only in the 'gvar' table). An intermediate-region
150 tuple variation table additionally has start and end n-tuples that also get used in the
151 interpolation process; these are always represented using embedded tuple records.

152 Also note that the privatePointNumbers flag is independent of the
153 SHARED_POINT_NUMBERS flag in the tupleVariationCount field of the
154 GlyphVariationData or 'cvar' header. A GlyphVariationData or 'cvar' table may have
155 shared point number data used by multiple tuple variation tables, but any given tuple
156 variation table may have private point number data that it uses instead.

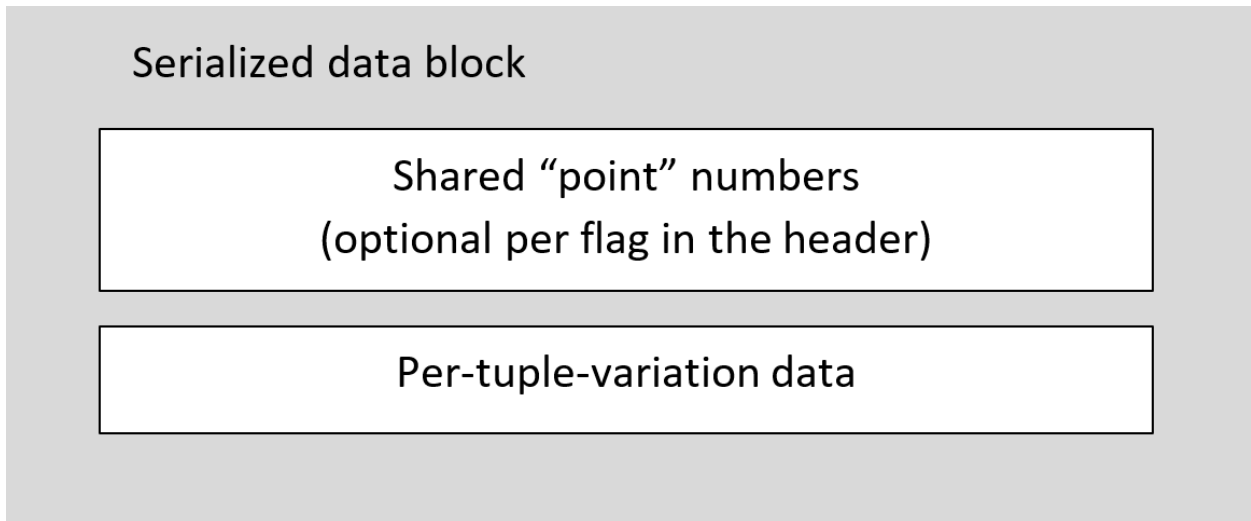
157 As noted, the size of tuple variation headers is variable. The next TupleVariationHeader
158 can be calculated as follows:

```
159 const TupleVariationHeader*  
160 NextHeader( const TupleVariationHeader* currentHeader, int axisCount )  
161 {  
162     int bump = 2 * sizeof( uint16 );  
163     int tupleIndex = currentHeader->tupleIndex;  
164     if ( tupleIndex & embeddedPeakTuple )  
165         bump += axisCount * sizeof( F2DOT14 );  
166     if ( tupleIndex & intermediateRegion )  
167         bump += 2 * axisCount * sizeof( F2DOT14 );  
168     return (const TupleVariationHeader*)((char*)currentHeader + bump);  
169 }
```

170 Serialized Data

171 After the GlyphVariationData or 'cvar' header (including the TupleVariationHeader array)
172 is a block of serialized data. The offset to this block of data is provided in the header.

173 The serialized data block begins with shared "point" number data, followed by the
174 variation data for the tuple variation tables. The shared point number data is optional: it
175 is present if the corresponding flag is set in the tupleVariationCount field of the header.
176 If present, the shared number data is represented as packed point numbers, described
177 below.

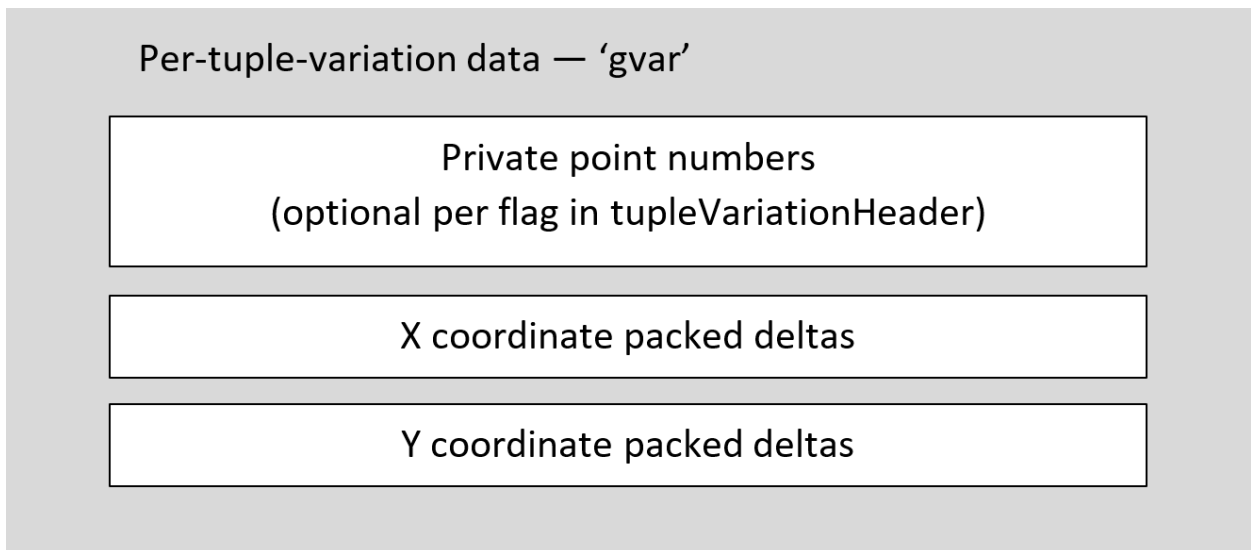


178
179 Figure: Organization of serialized data

180 The remaining data contains runs of data specific to individual tuple variation tables, in
 181 order of the tuple variation headers. Each TupleVariationHeader indicates the data size
 182 for the corresponding run of data for that tuple variation table.

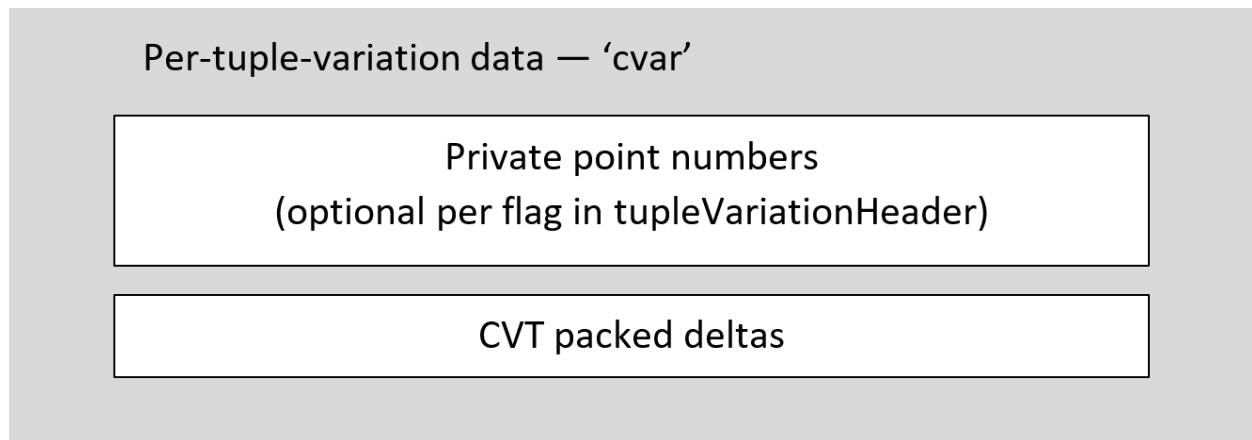
183 The per-tuple-variation-table data optionally begins with private “point” numbers,
 184 present if the privatePointNumbers flag is set in the tupleIndex field of the
 185 TupleVariationHeader. Private point numbers are represented as packed point numbers,
 186 described below.

187 After the private point number data (if present), the tuple variation data will include
 188 packed delta data. The format for packed deltas is given below. Within the 'gvar' table,
 189 there are packed deltas for X coordinates, followed by packed deltas for Y coordinates.



190
191 Figure: Organization 'gvar' per-tuple variation data

192 Within the 'cvar' table, there is one set of packed deltas.



193

194 Figure: Organization 'cvar' per-tuple variation data

195 The data size indicated in the TupleVariationHeader includes the size of the private
196 point number data, if present, plus the size of the packed deltas.

197 Packed "Point" Numbers

198 Tuple variation data specify deltas to be applied to specific items: X and Y coordinates
199 for glyph outline points within the 'gvar' table, and CVT values in the 'cvar' table. For a
200 given glyph, deltas may be provided for any or all of a glyph's points, including
201 "phantom" points generated within the rasterizer that represent glyph side bearing
202 points. (See the chapter [Instructing TrueType Glyphs](#) for more background on phantom
203 points.) Similarly, within the 'cvar' table, deltas may be provided for any or all CVTs. The
204 set of glyph points or CVTs for which deltas are provided is specified by packed point
205 numbers.

206 *Note:* If a glyph is a composite glyph, then "point" numbers are component indices for
207 the components that make up the composite glyph. See the 'gvar' table chapter for
208 complete details. Likewise, in the context of the 'cvar' table, "point" numbers are
209 indices for CVT entries.

210 *Note:* Within the 'gvar' table, if deltas are not provided explicitly for some points, then
211 inferred delta values may need to be calculated — see the 'gvar' table chapter for
212 details. This does not apply to the 'cvar' table, however: if deltas are not provided for
213 some CVT values, then no adjustments are made to those CVTs in connection to the
214 given tuple variation table.

215 Packed point numbers are stored as a count followed by one or more runs of point
216 number data.

217 The count may be stored in one or two bytes. After reading the first byte, the need for a
218 second byte can be determined. The count bytes are processed as follows:

- 219 • If the first byte is 0, then a second count byte is not used. This value has a special
220 meaning: the tuple variation data provides deltas for all glyph points (including
221 the “phantom” points), or for all CVTs.
- 222 • If the first byte is non-zero and the high bit is clear (value is 1 to 127), then a
223 second count byte is not used. The point count is equal to the value of the first
224 byte.
- 225 • If the high bit of the first byte is set, then a second byte is used. The count is read
226 from interpreting the two bytes as a big-endian uint16 value with the high-order
227 bit masked out.

228 Thus, if the count fits in 7 bits, it is stored in a single byte, with the value 0 having a
229 special interpretation. If the count does not fit in 7 bits, then the count is stored in the
230 first two bytes with the high bit of the first byte set as a flag that is not part of the count
231 — the count uses 15 bits.

232 For example, a count of 0x00 indicates that deltas are provided for all point numbers /
233 all CVTs, with no additional point number data required; a count of 0x32 indicates that
234 there are a total of 50 point numbers specified; a count of 0x81 0x22 indicates that there
235 are a total of 290 (= 0x0122) point numbers specified.

236 Point number data runs are given after the count. Each data run begins with a control
237 byte that specifies the number of point numbers defined in the run, and a flag bit
238 indicating the format of the run data. The control byte’s high bit specifies whether the
239 run is represented in 8-bit or 16-bit values. The low 7 bits specify the number of
240 elements in the run minus 1. The format of the control byte is as follows:

Mask	Name	Description
0x80	POINTS_ARE_WORDS	Flag indicating the data type used for point numbers in this run. If set, the point numbers are stored as unsigned 16-bit values (uint16); if clear, the point numbers are stored as unsigned bytes (uint8).
0x7F	POINT_RUN_COUNT_MASK	Mask for the low 7 bits of the control byte to give the number of point number elements, minus 1.

241 For example, a control byte of 0x02 indicates that the run has three elements
242 represented as uint8 values; a control byte of 0xD4 indicates that the run has $0x54 + 1 =$
243 85 elements represented as uint16 values.

244 In the first point run, the first point number is represented directly (that is, as a
245 difference from zero). Each subsequent point number in that run is stored as the
246 difference between it and the previous point number. In subsequent runs, all elements,
247 including the first, represent a difference from the last point number.

248 Since the values in the packed data are all unsigned, point numbers will be given in
249 increasing order. Since the packed representation can include zero values, it is possible
250 for a given point number to be repeated in the derived point number list. In that case,
251 there will be multiple delta values in the deltas data associated with that point number.
252 All of these deltas must be applied cumulatively to the given point.

253 Packed Deltas

254 Tuple variation data specify deltas to be applied to glyph point coordinates or to CVT
255 values. As in the case of point number data, deltas are stored in a packed format.

256 Packed delta data does not include the total number of delta values within the data.
257 Logically, there are deltas for every point number or CVT index specified in the point-
258 number data. Thus, the count of logical deltas is equal to the count of point numbers
259 specified for that tuple variation table. But since the deltas are represented in a packed
260 format, the actual count of stored values is typically less than the logical count. The data
261 is read until the expected logic count of deltas is obtained.

262 **Note:** In the 'gvar' table, there will be two logical deltas for each point number: one
263 that applies to the X coordinate, and one that applies to the Y coordinate. Therefore,
264 the total logical delta count is two times the point number count. The packed deltas
265 are arranged with the deltas for X coordinates first, followed by the deltas for Y
266 coordinates.

267 Packed deltas are stored as a series of runs. Each delta run consists of a control byte
268 followed by the actual delta values of that run. The control byte is a packed value with
269 flags in the high two bits and a count in the low six bits. The flags specify the data size
270 of the delta values in the run. The format of the control byte is as follows:

Mask	Name	Description
0x80	DELTAS_ARE_ZERO	Flag indicating that this run contains no data (no explicit delta values are stored), and that the deltas for this run are all zero.
0x40	DELTAS_ARE_WORDS	Flag indicating the data type for delta values in the run. If set, the run contains 16-bit signed deltas (int16); if clear, the run contains 8-bit signed deltas (int8).
0x3F	DELTA_RUN_COUNT_MASK	Mask for the low 6 bits to provide the number of delta values in the run, minus one.

271 For example, a control byte of 0x03 indicates that there are four 8-bit signed delta
272 values following the control byte; a control byte of 0x40 indicates that there is one 16-
273 bit signed delta value following the control byte; a control byte of 0x94 indicates that
274 there is no additional data for this run, and that the run represents a sequence of 0x14 +
275 1 = 21 deltas equal to zero.

276 ■ 03 0A 97 00 C6 87 41 10 22 FB 34

277 This data has three runs: a run of four 8-bit values, a run interpreted as eight zeroes, and
278 a run of two 16-bit values:

279 ■ Run 1: 03 0A 97 00 C6

280 ■ Run 2: 87

281 ■ Run 3: 41 10 22 FB 34

282 This packed data would represent the following logical sequence of delta values:

283 ■ 10, -105, 0, -58, 0, 0, 0, 0, 0, 0, 0, 0, 0, 4130, -1228

284 Processing Tuple Variation Store Data

285 When a variation instance has been selected, an application needs to process the
286 variation store data to derive interpolated values for that instance — interpolated grid
287 coordinates for outline points, or interpolated CVT values. In the case of the 'gvar' table,
288 this will be done glyph-by-glyph as needed. The application can process the
289 TupleVariationHeaders to filter the tuple variation tables that are applicable for the
290 current instance, or to calculate a scalar for each tuple variation table directly. Scalars

291 can then be applied to deltas in each tuple variation table, and the net adjustments
292 applied to the target items.

293 *Note:* In the 'cvar' table, there is a logical delta for each CVT index given in the packed
294 point number data. In the 'gvar' table, there are two logical deltas for each point
295 number: one for the point's X coordinate, and one for the Y coordinate. The delta data
296 is organized with all of the deltas for X coordinates first, followed by deltas for Y
297 coordinates.

298 *Note:* In the 'gvar' table, if the data for a given glyph lists point numbers for some
299 points in a contour but not others, then delta values for the omitted point numbers
300 must be inferred. See the ['gvar' table](#) chapter for details.

301 For details on determining applicability of a given tuple variation table, and on
302 calculation of scalars and net adjustments to target items, see the chapter [OpenType
303 Font Variations Overview](#).

304 Because point number and delta data are stored in a packed representation, the data
305 must be processed from the start to determine the presence of any particular point
306 number, or to retrieve the delta for a particular item. For this reason, the format is best
307 suited to processing all the data in a given tuple variation table at once rather than
308 processing data for individual target items. In the case of glyph outlines, this is
309 reasonable since there is no common application scenario for interpolating an adjusted
310 position of a single outline point.

311 The "phantom" points, which provide side-bearing and advance width information, are a
312 possible exception to that generalization, however. (See the chapter, [Instructing
313 TrueType Glyphs](#) for more background on phantom points.) In particular, some text-
314 layout operations require glyph metrics (advance widths or side bearings) without
315 necessarily requiring glyph outline data. Yet the tuple variation store formats used in the
316 'gvar' table require that interpolated outlines be computed to obtain the interpolated
317 glyph metrics. The [HVAR](#) table and [VVAR](#) table provide an alternate way to represent
318 horizontal and vertical glyph metric variation data, and these use the item variation
319 store format which is specifically designed to be suitable for processing data for
320 particular target items.

321 Differences Between 'gvar' and 'cvar' Tables

322 The following is a summary of key differences between tuple variation stores in the
323 'gvar' and 'cvar' tables.

- 324
- 325
- 326
- 327
- 328
- 329
- 330
- 331
- 332
- 333
- 334
- 335
- 336
- 337
- 338
- 339
- 340
- 341
- 342
- 343
- 344
- 345
- 346
- 347
- 348
- 349
- 350
- 351
- 352
- The 'gvar' table is a parent table for tuple variation stores, and contains one tuple variation store (the glyph variation data table) for each glyph ID. In contrast, the entire 'cvar' table is comprised of a single, slightly extended (with version fields) tuple variation store.
 - Because the 'gvar' table contains multiple tuple variation stores, sharing of data between tuple variation stores is possible, and is used for shared tuple records. Because the 'cvar' table has a single tuple variation store, no possibility of shared data arises.
 - The tupleIndex field of TupleVariationHeader structures within a tuple variation store includes a flag that indicates whether the structure instance includes an embedded peak tuple record. In the 'gvar' table, this is optional. In the 'cvar' table, a peak tuple record is mandatory.
 - The serialized data includes packed “point” numbers. In the 'gvar' table, these refer to glyph contour point numbers or, in the case of a composite glyph, to component indices. In the context of the 'cvar' table, these are indices for CVT entries.
 - In the 'gvar' table, point numbers cover the points or components defined in a 'glyf' entry plus four additional “phantom” points that represent the glyph’s horizontal and vertical advance and side bearings. (See the chapter, [Instructing TrueType Glyphs](#) for more background on phantom points.) The last four point numbers for any glyph, including composite glyphs, are for the phantom points.
 - In the 'gvar' table, if deltas are not provided for some points and the point indices are not represented in the point number data, then interpolated deltas for those points will in some cases be inferred. This is not done in the 'cvar' table, however.
 - In the 'gvar' table, the serialized data for a given region has two logical deltas for each point number: one for the X coordinate, and one for the Y coordinate. Hence the total number of deltas is twice the count of control points. In the 'cvar' table, however, there is only one delta for each point number.

353 Item Variation Store

354 Item variation stores are used for most variation data other than that used for TrueType
355 glyph outlines, including the variation data in [MVAR](#), [HVAR](#), [VVAR](#), [BASE](#) and [GDEF](#)
356 tables.

357 **Note:** For CFF2 glyph outlines, delta values are interleaved directly within the glyph
358 outline description in the CFF2 table. The sets of regions which are associated with the
359 delta sets are defined in an item variation store, contained as a subtable within the
360 CFF2 table. See the [CFF2](#) chapter for additional details.

361 The item variation store formats organize sets of variation data into groupings by the
362 target items. This makes the formats well-suited to computing interpolated instance
363 values for individual font data items. This is useful for certain text layout operations in
364 which only certain data items are required, such as the advance widths of specific glyphs
365 or anchor positions used in specific GPOS lookup tables.

366 The different tables that use item variation stores have their own top-level formats. Each
367 will include an offset to an itemVariationStore table, containing the variation data. This
368 chapter describes the shared formats: the itemVariationStore and its component
369 structures.

370 Associating Target Items to Variation Data

371 Variation data is comprised of delta adjustment values that apply to particular target
372 items. Some mechanism is needed to associate delta values with target items. In the
373 item variation store, a block of delta values has an implicit delta-set index, and separate
374 data outside the item variation store is provided that indicates the delta-set index
375 associated with a given target item. Depending on the parent table in which an item
376 variation store is used, different means are used to provide these associations:

- 377 • In the MVAR table, an array of records identifies target data items in various
378 other tables, along with the delta-set index for each respective item.
- 379 • In the HVAR and VVAR tables, the target data items are glyph metric arrays in the
380 'hmtx' and 'vmtx' tables. Subtables in the HVAR and VVAR tables provide the
381 mapping between the target data items and delta-set indices.
- 382 • For the BASE, GDEF, GPOS, and JSTF tables, a target data item is associated with a
383 delta-set index using a related [VariationIndex table](#) within the same subtable that
384 contains the target item.
- 385 • In the COLR table, target data items are specified in structures that combine a
386 basic data type, such FWORD, with a delta-set index.

387 The structures used in the COLR table currently are used only in that table but may be
388 used in other tables in future versions, and so are defined here as common formats.
389 Structures are defined to wrap the FWORD, UFWORD, F2DOT14 and Fixed basic types.

390 Note: as described below, each delta-set index is represented as two index components,
391 an *outer* index and an *inner* index, corresponding to a two-level organizational
392 hierarchy. This is described in detail below.

393 **VarFWord**

394 The FWORD type is used to represent coordinates in the glyph design grid. The
395 VarFWord record is used to represent a coordinate that can be variable.

Type	Name	Description
FWORD	coordinate	
uint16	varOuterIndex	
uint16	varInnerIndex	

396 **VarUFWord**

397 The UFWord type is used to represent distances in the glyph design grid. The
398 VarUFWord record is used to represent a distance that can be variable.

Type	Name	Description
UFWORD	distance	
uint16	varOuterIndex	
uint16	varInnerIndex	

399 **VarF2Dot14**

400 The F2DOT14 type is typically used to represent values that are inherently limited to a
401 range of [-1, 1], or a range of [0, 1]. The VarF2Dot14 record is used to represent such a
402 value that can be variable.

Type	Name	Description
F2Dot14	value	
uint16	varOuterIndex	
uint16	varInnerIndex	

403 In general, variation deltas are (logically) signed 16-bit integers, and in most cases, they
404 are applied to signed 16-bit values (FWORDS) or unsigned 16-bit values (UFWORDs).
405 When scaled deltas are applied to F2DOT14 values, the F2DOT14 value is treated like a
406 16-bit integer. (In this sense, the delta and the F2DOT14 value can be viewed as an
407 integral numerator for 1/16384ths.)

408 If the context in which the VarF2Dot14 is used constrains the valid range for the default
409 value, then any variations by applying deltas are clipped to that range.

410 **VarFixed**

411 The Fixed type is intended for floating values, such as variation-space coordinates. The
412 VarFixed record is used to represent such a value that can be variable.

Type	Name	Description
Fixed	value	
uint16	varOuterIndex	
uint16	varInnerIndex	

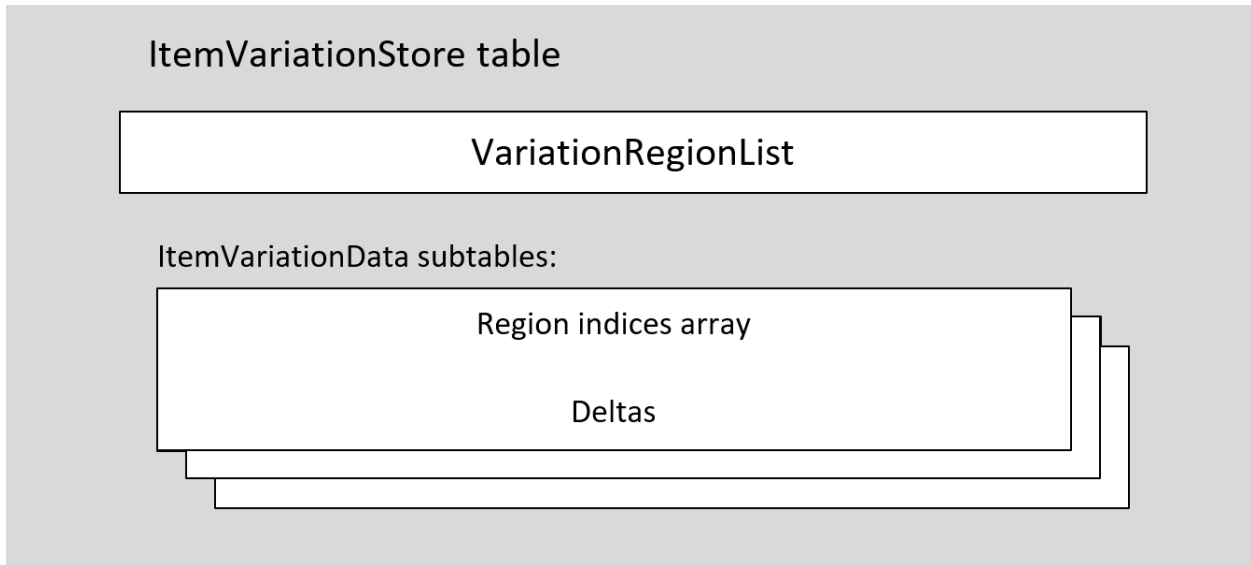
413 While in most cases deltas are applied to 16-bit types, Fixed is a 32-bit (16.16) type and
414 requires 32-bit deltas. The DeltaSet record used in the ItemVariationData subtable
415 format can accommodate deltas that are, logically, either 16-bit or 32-bit. See the
416 description of the [ItemVariationData subtable](#), below, for details.

417 When scaled deltas are applied to Fixed values, the Fixed value is treated like a 32-bit
418 integer. (In this sense, the delta and the Fixed value can be viewed as an integral
419 numerator of 1/65536ths.)

420 Variation Data

421 The ItemVariationStore table includes a variation region list, which defines the different
422 regions of the font's variation space for which variation data is defined. It also includes a
423 set of itemVariationData subtables, each of which provides a portion of the total
424 variation data. Each subtable is associated with some subset of the defined regions, and
425 will include deltas used for one or more target items. Conceptually, the deltas form a
426 two-dimensional array, with delta-set rows that include a delta for each of the regions
427 referenced by that subtable. From this perspective, the table columns correspond to
428 regions.

429 The following figure illustrates the overall structure of the ItemVariationStore stable.



430
431 Figure: High-level organization of ItemVariationStore table

432 Note that multiple subtables are necessary only if the number of distinct delta-set data
433 exceeds 65,536. Multiple subtables may also be used, however, to provide more
434 compact data representation. There are different ways that the delta data can be made
435 more compact.

436 First, deltas with a value of zero have no impact on their target items. If there are several
437 delta-set rows that have a zero delta for the same region, then those rows could be
438 moved into a subtable that does not reference that region. As a result, there will be
439 fewer delta values in each row, making the size of data for those rows smaller.

440 Also, some delta values require 16-bit representations, but some require only 8 bits. For
441 a given subtable, deltas in each row correspond, in order, to the regions that are
442 referenced, but the ordering of regions has no effect. Hence, regions and corresponding
443 deltas within each row can be reordered. Thus, regions that require 16-bit delta
444 representations can be ordered together. The itemVariationData format specifies a
445 count of regions (columns) for which a 16-bit delta representation is used, with the
446 remaining deltas in each row using 8 bits. By reordering columns, the size required for a
447 given delta-set row can potentially be reduced. If a set of rows have similar
448 requirements in regard to which columns have deltas requiring 16-bit versus 8-bit
449 representations, then those rows can be moved into a subtable with a column order that
450 allows a maximal number of deltas using 8-bit rather than 16-bit representations.

451 Note that there is minimal overhead for each subtable: 10 bytes (6 bytes in the subtable
452 header and 4 bytes for the offset in the parent table) plus 2 bytes for each region that is
453 referenced.

454 A complete delta-set index involves an outer-level index into the ItemVariationData
455 subtable array, plus an inner-level index to a delta-set row within that subtable. A special
456 meaning is assigned to a delta-set index 0xFFFF/0xFFFF (that is, outer-level and inner-
457 level portions are both 0xFFFF): this is used to indicate that there is no variation data for
458 a given item. Functionally, this would be equivalent to referencing delta-set data
459 consisting of only deltas of 0 for all regions.

460 As noted above, delta-set indices are stored outside the variation store. Different parent
461 tables that use an item variation store will store indices in different ways, and may utilize
462 different schemes for how to represent the indices in an efficient manner. For example,
463 the HVAR and VVAR tables allow the outer and inner indices to be combined into one-
464 byte, two-byte, three-byte or four-byte representations depending on the indexing
465 requirements of the variation store. For larger sets of variation data, such as may be
466 needed for HVAR or VVAR tables, optimization of the indices data as well as the delta
467 data may have a significant impact on overall size. Optimizing compilers may need to
468 consider the impact on representation of indices in tandem as it optimizes the item
469 variation store to achieve the best overall results.

470 Variation Regions

471 As noted above, variation data is comprised of delta adjustment values that have effect
472 over particular regions within the font's variation space. In a tuple variation store
473 (described earlier in this chapter), the deltas are organized into groupings by region of
474 applicability, with each grouping associated with a given region. In contrast, the item
475 variation store format organizes deltas into groupings by the target items to which they
476 apply, with each grouping having deltas for several regions. Accordingly, the item
477 variation store uses different formats for describing the regions in which a set of deltas
478 apply.

479 For a given item variation store, a set of regions is specified using a VariationRegionList.

480 *VariationRegionList:*

Type	Name	Description
uint16	axisCount	The number of variation axes for this font. This must be the same number as axisCount in the 'fvar' table.
uint16	regionCount	The number of variation region tables in the variation region list. Must be less than 32,768.

VariationRegion	variationRegions[regionCount]	Array of variation regions.
-----------------	-------------------------------	-----------------------------

481 The high-order bit of the regionCount field is reserved for future use, and must be
482 cleared.

483 The regions can be in any order. The regions are defined using an array of
484 RegionAxisCoordinates records, one for each axis defined in the 'fvar' table:

485 *VariationRegion record:*

Type	Name	Description
RegionAxisCoordinates	regionAxes[axisCount]	Array of region axis coordinates records, in the order of axes given in the 'fvar' table.

486 Each RegionAxisCoordinates record provides coordinate values for a region along a
487 single axis:

488 *RegionAxisCoordinates record:*

Type	Name	Description
F2DOT14	startCoord	The region start coordinate value for the current axis.
F2DOT14	peakCoord	The region peak coordinate value for the current axis.
F2DOT14	endCoord	The region end coordinate value for the current axis.

489 The three values must all be within the range -1.0 to +1.0. startCoord must be less than
490 or equal to peakCoord, and peakCoord must be less than or equal to endCoord. The
491 three values must be either all non-positive or all non-negative with one possible
492 exception: if peakCoord is zero, then startCoord can be negative or 0 while endCoord
493 can be positive or zero.

494 *Note:* The following guidelines are used for setting the three values in different
495 scenarios:

- 496 • In the case of a non-intermediate region for which the given axis should factor
497 into the scalar calculation for the region, either startCoord and peakCoord are
498 set to a negative value (typically, -1.0) and endCoord is set to zero, or
499 startCoord is set to zero and peakCoord and endCoord are set to a positive
500 value (typically +1.0).

- In the case of an intermediate region for which the given axis should factor into the scalar calculation for the region, startCoord, peakCoord and endCoord are all set to non-positive values or are all set to non-negative values.
- If the given axis should not factor into the scalar calculation for a region, then this is achieved by setting peakCoord to zero. In this case, startCoord can be any non-positive value, and endCoord can be any non-negative value. It is recommended either that all three be set to zero, or that startCoord be set to -1.0 and endCoord be set to +1.0.

The full algorithm for interpolation of instance values is given in the chapter, [OpenType Font Variations Overview](#). The logical algorithm involves computing per-axis scalar values for a given region and a given instance. The per-axis scalars for a region are then combined to yield an overall scalar for the region that is then applied to delta adjustment values. Given a selected variation instance, a per-axis scalar can be calculated for each RegionAxisCoordinates record. The overall scalar for a region can be calculated by combining the per-axis scalars for that region.

Item Variation Store Header and Item Variation Data Subtables

The item variation store table has a header with the following structure.

ItemVariationStore table:

Type	Name	Description
uint16	format	Format — set to 1
Offset32	variationRegionListOffset	Offset in bytes from the start of the item variation store to the variation region list.
uint16	itemVariationDataCount	The number of item variation data subtables.
Offset32	itemVariationDataOffsets[itemVariationDataCount]	Offsets in bytes from the start of the item variation store to each item variation data subtable.

The item variation store includes an offset to a variation region list and an array of offsets to item variation data subtables.

521 *Note:* Indices into the itemVariationDataOffsets array are stored in parent tables as
 522 delta-set “outer” indices with each such index having a corresponding “inner” index. If
 523 the outer index points to a NULL offset, then any inner index will be invalid. The
 524 itemVariationDataOffsets array should not include any NULL offsets.

525 Each item variation data subtable includes deltas for some number of items, and some
 526 subset of regions. The regions are indicated by an array of indices into the variation
 527 region list.

528 *ItemVariationData subtable:*

Type	Name	Description
uint16	itemCount	The number of delta sets for distinct items.
uint16	wordDeltaCount	A packed field: the high bit is a flag—see details below.
uint16	regionIndexCount	The number of variation regions referenced.
uint16	regionIndexes[regionIndexCount]	Array of indices into the variation region list for the regions referenced by this item variation data table.
DeltaSet	deltaSets[itemCount]	Delta-set rows.

529 The wordDeltaCount field contains a packed value that includes a flag and a “word”
 530 delta count. The format of this value is as follows:

Mask	Name	Description
0x8000	LONG_WORDS	Flag indicating that “word” deltas are long (int32)
0x7FFF	WORD_DELTA_COUNT_MASK	Count of “word” deltas

531 The representation of delta values uses a mix of long types (“words”) and short types. If
 532 the LONG_WORDS flag is set, deltas are represented using a mix of int32 and int16
 533 values. This representation is only used for deltas that are to be applied to data items of
 534 Fixed or 32-bit integer types. If the flag is not set, deltas are presented using a mix of
 535 int16 and int8 values. See the description of the DeltaSet record below for additional
 536 details.

537 The count value indicated by WORD_DELTA_COUNT_MASK is a count of the number of
 538 deltas that use the long (“word”) representation, and must be less than or equal to
 539 regionIndexCount.

540 The deltaSets array represents a logical two-dimensional table of delta values with
541 itemCount rows and regionIndexCount columns. Rows in the table provide sets of deltas
542 for particular target items, and columns correspond to regions of the variation space.
543 Each DeltaSet record in the array represents one row of the delta-value table — one
544 delta set.

545 *DeltaSet record:*

Type	Name	Description
int16 and int8 <i>or</i> int32 and int16	deltaData[regionIndexCount]	Variation delta values.

546 Logically, each DeltaSet record has regionIndexCount number of elements. The
547 elements are represented using long and short types, as described above. These are
548 either int16 and int8, or int32 and int16, according to whether the LONG_WORDS flag
549 was set. The delta array has a sequence of deltas using the long type followed by
550 sequence of deltas using the short type. The count of deltas using the long type is
551 derived using WORD_DELTA_COUNT_MASK. The remaining elements use the short type.
552 The length of the data for each row, in bytes, is regionIndexCount + (wordDeltaCount
553 && WORD_DELTA_COUNT_MASK) if the LONG_WORDS flag is not set, or 2 x that
554 amount if the flag is set.

555 **Note:** Delta values are each represented directly. They are not packed as in the tuple
556 variation store.

557 Processing Item Variation Store Data

558 When a variation instance has been selected, an application needs to process the
559 variation store data associated with particular target items to derive interpolated values
560 for those items and that instance.

561 To compute the interpolated instance value for a given target item, the application first
562 obtains the delta-set index for that item. It uses the outer-level index portion to select
563 an item variation data subtable within the item variation store, and the inner-level index
564 portion to select a delta-set row within that subtable. The delta set contains one delta
565 for each region referenced by the subtable, in order of the region indices given in the
566 regionIndices array. The application uses the regionIndices array for that subtable to
567 identify applicable regions and to compute a scalar for each of these regions based on
568 the selected instance. Each of the scalars is then applied to the corresponding delta

569 within the delta set to derive a scaled adjustment. The scaled adjustments for the row
570 are then combined to obtain the overall adjustment for the item.

571 Complete details on the interpolation algorithm logic are provided in the chapter,
572 [OpenType Font Variations Overview](#).

573 For a given variation instance, an application will often need to interpolate values for
574 several items that may use deltas in different item variation data subtables. The
575 subtables will reference region definitions in the shared variation region list. When the
576 instance has been selected, applications can pre-compute and cache a scalar for that
577 instance for each region in the region list. Then when processing different target items,
578 the cached scalar array can be used without needing to re-compute region scalars for
579 each target item.